

End-to-end evaluation pipeline (pooled 32-slot scoring)

Panoramic radiograph

Tufts: 32 Universal slots (1-32)
Missing vs Present

DENTEX: 32 FDI slots (11-48)
Abnormal vs No annotated finding

VLM · LLaVA-1.5-7B · LLaVA-Med · HuatuoGPT-Vision · DentVLM
Prompting strategy · zero-shot · few-shot · chain-of-thought

Parser → tooth-level predictions
Abnormality from parsed FDI/Universal findings (not from Total-count lines)

One tooth, two findings
FDI 46: Caries + Periapical
→ yes in both OvR classifiers

Table A / Tufts binary
Pooled TP/FP/FN/TN
Headline: Macro F1
Also: Acc, BalAcc

DENTEX Table A
Abnormal vs no finding
Headline: Macro F1
Also: Acc, BalAcc

DENTEX Table B (multi-label)
4 independent one-vs-rest
Impacted · Caries
Deep Caries · Periapical
Headline: Macro-4 F1